

Website Fingerprinting using Side-Channel Attacks: Implementation and Analysis

Rafiqul Islam Rayan
Student ID: 2005062
Department of Computer Science and Engineering
Bangladesh University of Engineering and Technology

June 20, 2025

Contents

1	Introduction	3
1.1	Motivation	3
1.2	Objectives	3
2	Methodology	3
2.1	Overall Architecture	3
2.2	Technical Stack	3
3	Phase 1: Latency Measurement	4
3.1	Implementation	4
3.1.1	Key Parameters	4
3.1.2	Algorithm	4
3.2	Results	4
3.3	Analysis	4
4	Phase 2: Cache Sweep Implementation	5
4.1	Design	5
4.1.1	Parameters	5
4.1.2	Algorithm	5
4.2	Challenges Faced	6
4.3	Solutions Implemented	6
5	Phase 3: Dataset Generation	6
5.1	Data Collection Process	6
5.1.1	Target Websites	6
5.1.2	Dataset Structure	6
5.2	Data Characteristics	6
5.3	Data Quality Issues	6
6	Phase 4: Machine Learning Model Training	7
6.1	Model Architecture	7
6.1.1	Basic CNN Model	7
6.1.2	Complex CNN Model	7
6.2	Training Configuration	7

6.3	Model Performance	7
6.4	Challenges in Training	7
7	Phase 5: Real-time Detection System	8
7.1	System Architecture	8
7.1.1	Components	8
7.2	Implementation Details	8
7.3	Real-time Challenges	8
8	Detailed Analysis and Findings	9
8.1	Cache Latency Characterization	9
8.2	Timing Trace Analysis	9
8.3	Model Performance Analysis	9
8.3.1	Basic CNN Performance	9
8.3.2	Complex CNN Performance	9
8.4	Real-time System Performance	10
9	Challenges and Technical Difficulties	10
9.1	Browser Security Constraints	10
9.2	Data Collection Challenges	10
9.3	Machine Learning Limitations	10
10	Advanced Techniques and Future Improvements	10
10.1	Multi-Channel Side-Channel Fusion	10
10.2	Evasion Techniques	11
10.3	Improved Data Collection	11
11	Results and Discussion	11
11.1	Key Findings	11
11.2	Attack Accuracy	11
11.3	Limitations Identified	12
12	Security Implications	12
12.1	Privacy Concerns	12
12.2	Potential Countermeasures	12
13	Conclusion	12

1 Introduction

Website fingerprinting is a privacy attack that allows an adversary to infer which websites a user is visiting by analyzing side-channel information such as network traffic patterns, timing characteristics, or cache behavior. This project implements a browser-based side-channel attack using JavaScript to collect timing data and machine learning techniques to classify websites.

1.1 Motivation

Understanding side-channel vulnerabilities is crucial for both security researchers and practitioners. This project demonstrates:

- How seemingly innocuous timing information can leak sensitive browsing patterns
- The effectiveness of machine learning in pattern recognition for security applications
- Real-world challenges in implementing side-channel attacks in modern browsers
- Potential countermeasures and their limitations

1.2 Objectives

1. Implement cache timing measurements in JavaScript
2. Develop data collection mechanisms for website fingerprinting
3. Generate a comprehensive dataset of timing traces
4. Train machine learning models for website classification
5. Deploy a real-time detection system

2 Methodology

2.1 Overall Architecture

The project follows a five-phase approach:

1. **Phase 1:** Latency Measurement - Characterizing cache access patterns
2. **Phase 2:** Cache Sweep Implementation - Collecting timing traces
3. **Phase 3:** Dataset Generation - Creating labeled training data
4. **Phase 4:** Model Training - Developing classification algorithms
5. **Phase 5:** Real-time Detection - Deploying the attack system

2.2 Technical Stack

- **Frontend:** HTML5, JavaScript (Web Workers)
- **Backend:** Python, Flask
- **Machine Learning:** PyTorch, scikit-learn
- **Browser Automation:** Selenium WebDriver
- **Data Processing:** NumPy, pandas

3 Phase 1: Latency Measurement

3.1 Implementation

The first phase involved implementing a cache latency measurement system using JavaScript Web Workers to characterize memory access patterns.

3.1.1 Key Parameters

- **Cache Line Size:** 64 bytes (determined via system configuration)
- **Measurement Iterations:** 10 repetitions per test
- **Buffer Sizes:** 1 to 10,000,000 cache lines
- **Timing Precision:** `performance.now()` API

3.1.2 Algorithm

```

1 function readNlines(n) {
2   const buffer = new ArrayBuffer(n * LINESIZE);
3   const times = [];
4
5   for (let i = 0; i < 10; i++) {
6     const start = performance.now();
7     for (let j = 0; j < n; j++) {
8       const offset = j * LINESIZE;
9       const line = buffer[offset]; // Cache access
10    }
11    const end = performance.now();
12    times.push(end - start);
13  }
14
15  return median(times);
16 }
```

Listing 1: Cache Latency Measurement Algorithm

3.2 Results

Table 1: Cache Access Latency Measurements

Cache Lines (N)	Median Access Latency (ms)
1	0.00
10	0.00
100	0.00
1,000	0.00
10,000	0.00
100,000	0.00-0.1
1,000,000	0.5-1.00
10,000,000	5.00

3.3 Analysis

The results demonstrate clear cache hierarchy effects:

- **L1/L2 Cache:** Negligible latency for smaller buffer sizes ($< 100,000$ lines)
- **L3 Cache:** Slight latency increase at 1,000,000 lines
- **Main Memory:** Significant latency (5ms) for large buffers (10,000,000 lines)

This validates our understanding of the memory hierarchy and confirms that timing measurements can distinguish between different cache levels.

4 Phase 2: Cache Sweep Implementation

4.1 Design

The cache sweep phase implements continuous memory access patterns to collect timing traces that reflect website-specific cache behavior.

4.1.1 Parameters

- **LLC Size:** 24 MB (reduced for memory efficiency)
- **Total Time:** 10 seconds per trace
- **Sweep Period:** 10ms intervals
- **Expected Traces:** 1000 data points per website visit

4.1.2 Algorithm

```
1 function sweep(P) {
2   const buffer = new ArrayBuffer(LLCSIZE);
3   const view = new Uint8Array(buffer);
4   const numCacheLines = Math.floor(LLCSIZE / LINESIZE);
5   const numIntervals = Math.floor(TIME / P);
6   const traces = [];
7
8   for (let interval = 0; interval < numIntervals; interval++) {
9     let sweepCount = 0;
10    const startTime = performance.now();
11
12    while (performance.now() - startTime < P) {
13      for (let line = 0; line < numCacheLines; line++) {
14        const offset = line * LINESIZE;
15        const dummy = view[offset]; // Memory access
16      }
17      sweepCount++;
18    }
19    traces.push(sweepCount);
20  }
21  return traces;
22 }
```

Listing 2: Cache Sweep Implementation

4.2 Challenges Faced

1. **Browser Security Restrictions:** Modern browsers limit high-resolution timing
2. **Memory Constraints:** Large buffer allocations caused browser crashes
3. **Background Interference:** Other browser processes affected measurements
4. **Timing Precision:** JavaScript timing APIs have limited resolution

4.3 Solutions Implemented

- Reduced cache size from theoretical maximum to practical 24MB
- Used Web Workers for isolated execution environment
- Implemented progressive data collection with memory management
- Added error handling and data truncation for large datasets

5 Phase 3: Dataset Generation

5.1 Data Collection Process

The dataset generation phase involved collecting timing traces from multiple website visits to create labeled training data.

5.1.1 Target Websites

- <https://cse.buet.ac.bd/moodle/> - Academic platform
- <https://google.com> - Search engine
- <https://prothomalo.com> - News website

5.1.2 Dataset Structure

```

1 {
2   "website": "https://google.com",
3   "website_index": 1,
4   "trace_data": [44, 34, 42, 42, 43, ...] // 1000 values
5 }
```

Listing 3: Dataset Entry Format

5.2 Data Characteristics

Table 2: Dataset Statistics

Website	Samples	Trace Length	Mean Value	Std Dev
BUET Moodle	10	1000	41.54	0.13
Google	10	1000	41.44	0.32
Prothom Alo	10	1000	39.79	1.09

5.3 Data Quality Issues

1. **Incomplete Traces:** Some samples had fewer than 1000 data points

2. **Limited Diversity:** Small number of samples per website
3. **Temporal Correlation:** Traces collected sequentially may show temporal bias

6 Phase 4: Machine Learning Model Training

6.1 Model Architecture

Two convolutional neural network architectures were implemented and compared:

6.1.1 Basic CNN Model

- 2 Conv1D layers (32, 64 filters)
- MaxPooling and fully connected layers
- Dropout regularization (0.5)
- Parameters: 130K

6.1.2 Complex CNN Model

- 3 Conv1D layers (32, 64, 128 filters)
- Batch normalization
- Multiple dropout layers (0.5, 0.3)
- Parameters: 280K

6.2 Training Configuration

Table 3: Training Hyperparameters

Parameter	Value
Batch Size	64
Epochs	50
Learning Rate	1e-4
Optimizer	Adam
Loss Function	CrossEntropyLoss
Train/Test Split	80/20

6.3 Model Performance

Table 4: Model Comparison Results

Model	Parameters	Training Accuracy	Test Accuracy
Basic CNN	1,035,011	54.17%	50.00%
Complex CNN	4,161,859	100.00%	66.67%

6.4 Challenges in Training

1. **Limited Data:** Small dataset size led to potential overfitting

2. **Class Imbalance:** Unequal number of samples per website
3. **Feature Engineering:** Raw timing data required careful preprocessing
4. **Model Selection:** Balancing complexity vs. generalization

7 Phase 5: Real-time Detection System

7.1 System Architecture

The real-time detection system implements a Flask web application that can classify websites in real-time using the trained models.

7.1.1 Components

- **Flask Backend:** Model serving and prediction API
- **Selenium WebDriver:** Automated browser control
- **Web Interface:** Real-time prediction display
- **Data Collection:** Live timing measurement

7.2 Implementation Details

```
1 def collect_timing_trace():
2     driver = setup_chrome_driver()
3     timing_data = []
4
5     # Collect timing measurements
6     for i in range(INPUT_SIZE):
7         driver.get("data:text/html,<html>...</html>")
8         timing = driver.execute_script("""
9             return performance.timing.loadEventEnd -
10                performance.timing.navigationStart;
11         """)
12         timing_data.append(scaled_timing(timing))
13
14     # Make prediction
15     prediction = model.predict(timing_data)
16     return prediction
```

Listing 4: Real-time Prediction Pipeline

7.3 Real-time Challenges

1. **Timing Consistency:** Real-time data differed from training data
2. **Browser Overhead:** WebDriver introduced additional latency
3. **Prediction Bias:** Model showed preference for certain classes

8 Detailed Analysis and Findings

8.1 Cache Latency Characterization

The initial latency measurements revealed distinct patterns across different memory hierarchy levels:

1. **L1/L2 Cache Response:** For buffer sizes up to 100,000 cache lines, access latency remained consistently at 0.00ms, indicating efficient cache hits.
2. **L3 Cache Transition:** At 1,000,000 cache lines, a noticeable latency increase to 1.00ms suggested L3 cache utilization.
3. **Main Memory Access:** The dramatic jump to 5.00ms at 10,000,000 cache lines confirmed main memory access patterns.

This hierarchy validation was crucial for understanding the underlying timing behavior that enables website fingerprinting.

8.2 Timing Trace Analysis

Analysis of the collected timing traces revealed several important characteristics:

- **Website-Specific Patterns:** Different websites showed distinct timing signatures, with Prothom Alo exhibiting the most variable patterns (std dev: 1.09)
- **Consistency within Sites:** BUET Moodle showed the most consistent timing patterns (std dev: 0.13), suggesting stable cache behavior
- **Temporal Stability:** Traces collected from the same website showed good repeatability, validating the fingerprinting approach

8.3 Model Performance Analysis

The machine learning results provide insights into the viability of the attack:

8.3.1 Basic CNN Performance

- Achieved modest 50% test accuracy with 1M+ parameters
- Training accuracy (54.17%) close to test accuracy suggests good generalization
- Performance limited by small dataset size and simple architecture

8.3.2 Complex CNN Performance

- Perfect training accuracy (100%) with 4M+ parameters
- Test accuracy of 66.67% indicates overfitting due to limited data
- Despite overfitting, still outperformed the basic model on unseen data
- Higher capacity allowed better feature learning from timing patterns

8.4 Real-time System Performance

The deployed Flask application successfully demonstrated:

- Real-time data collection and model inference
- Web-based interface for attack demonstration
- Integration of browser automation with machine learning pipeline
- Practical challenges in deployment environment

9 Challenges and Technical Difficulties

9.1 Browser Security Constraints

Modern browsers implement several security measures that significantly impact timing attacks:

1. **Timer Resolution Reduction:** `performance.now()` precision is intentionally reduced to prevent timing attacks
2. **Site Isolation:** Process isolation between different origins limits cross-site timing analysis
3. **Memory Access Restrictions:** Direct memory manipulation is limited in JavaScript environments
4. **Background Process Interference:** Browser maintenance tasks introduce timing noise

9.2 Data Collection Challenges

Several technical difficulties emerged during data collection:

1. **Timing Consistency:** Variations in system load affected measurement consistency
2. **Cross-Platform Differences:** Timing behavior varied across different hardware configurations
3. **Network Dependencies:** Website loading times introduced additional timing variability

9.3 Machine Learning Limitations

The ML pipeline faced several constraints:

1. **Limited Dataset Size:** Only 30 complete traces available for training
2. **Overfitting Risk:** High-capacity models with limited data

10 Advanced Techniques and Future Improvements

10.1 Multi-Channel Side-Channel Fusion

Future work could incorporate multiple side-channels:

- **Branch Prediction Patterns:** Exploit branch predictor state changes
- **TLB Timing:** Translation Lookaside Buffer access patterns
- **Prefetcher Behavior:** Hardware prefetcher state analysis
- **Network Timing:** Combine with network-level fingerprinting

10.2 Evasion Techniques

To improve attack success against defenses:

- **Adaptive Timing:** Adjust measurement parameters based on detected countermeasures
- **Noise Reduction:** Advanced signal processing to filter defense mechanisms
- **Indirect Measurements:** Use secondary effects to infer cache state
- **Statistical Analysis:** Apply advanced statistical methods to detect subtle patterns

10.3 Improved Data Collection

Future improvements could include:

- **Longitudinal Studies:** Collect data over extended periods
- **Cross-Browser Analysis:** Test across multiple browser engines
- **Mobile Platforms:** Extend to mobile browser environments
- **Diverse Websites:** Expand to larger website corpus

11 Results and Discussion

11.1 Key Findings

1. **Feasibility:** Website fingerprinting using cache timing is feasible in JavaScript
2. **Browser Limitations:** Modern browsers impose significant constraints on timing precision
3. **ML Effectiveness:** CNN models can learn patterns from timing data
4. **Real-world Gaps:** Laboratory conditions differ significantly from real-world deployment

11.2 Attack Accuracy

The final system achieved the following performance metrics:

- Basic CNN Training accuracy: 54.17%
- Basic CNN Testing accuracy: 50.00%
- Complex CNN Training accuracy: 100.00%
- Complex CNN Testing accuracy: 66.67%
- Real-time detection: Functional with model-based predictions

The Complex CNN model showed clear signs of overfitting with perfect training accuracy but limited generalization. However, it still outperformed the Basic CNN on test data, achieving 66.67% accuracy on the 6-sample test set.

11.3 Limitations Identified

1. **Limited Website Coverage:** Only 3 websites in dataset
2. **Temporal Sensitivity:** Performance varies with system load
3. **Browser Dependency:** Results specific to Chrome/Chromium
4. **Network Variability:** Network conditions affect timing

12 Security Implications

12.1 Privacy Concerns

This work demonstrates that:

- Browsing patterns can be inferred from timing side-channels
- JavaScript-based attacks can bypass some browser security measures
- Machine learning amplifies the effectiveness of timing attacks
- Real-time monitoring is feasible with sufficient resources

12.2 Potential Countermeasures

1. **Timing Randomization:** Adding noise to timing APIs
2. **Cache Partitioning:** Isolating cache access between origins
3. **Resolution Reduction:** Limiting timer precision
4. **Resource Throttling:** Controlling memory allocation

13 Conclusion

This project successfully demonstrated the implementation of a comprehensive website fingerprinting attack using browser-based side-channel analysis. Through five distinct phases spanning from fundamental cache latency characterization to real-time deployment, we developed a complete pipeline capable of identifying websites based solely on timing measurements collected via JavaScript.